



# Manny Manifolds – Complete Documentation Export

## Manny Manifolds – Complete Documentation Export

**Exported:** October 22, 2025

**Contents:** Overview, Architecture, Design Laws & Principles, Mathematics & Algorithms, Cognitive Physics Extensions, Evaluation & Metrics

---

## SECTION 1: OVERVIEW

### 1.1 Introduction

*Purpose, vision, and core vocabulary*

#### Purpose of the System

Create a **living cognitive substrate** where data, computation, and learning are one evolving geometry.

The Manifold Engine is a self-adapting conversational geometry engine that learns through interaction. Threads (experiences) flow through a deformable manifold whose curvature encodes relationships and context—those flows reshape the manifold, creating continual learning, reasoning, and empathy.

#### One-Sentence Vision

**Manny Manifolds succeeds when it can learn, generalize, and explain its own thinking — a self-evolving geometric mind that humans can teach, understand, and trust.**

#### Quick Glossary

**Manifold:** The space of knowledge itself. Nodes = concepts, edges = relations. Curvature encodes relevance, context, and understanding. Geometry changes through interaction.

**Threads:** Active trajectories of experience or reasoning. Follow local geodesics; modify curvature along their path. Represent thought, exploration, conversation.

**Valence:** Scalar or multi-channel "energy" of experience (importance, affect, novelty). Determines how strongly threads reshape curvature. Emotional/attentional analogue.

**Lenses:** Contextual projections or coordinate systems. Define which dimensions are visible and how data are interpreted. Emergent from repeated thread traffic.

**Motifs:** Frequently traversed subpaths that become reusable skills or knowledge patterns. Cached for efficiency and analogical transfer.

**Drives:** Six-tier motivation hierarchy (stability, continuity, connection, competence, creativity, contribution) that arise from imbalance and guide energy allocation.

**Curvature:** The geometric property encoding learned associations. High curvature = strong, well-traveled connections. Learning = changing curvature.

**Plasticity:** Local Hebbian-style updates to curvature (online) plus consolidation ("sleep") for global re-projection and pruning (offline). Balance between learning and stability.

**Bicameral System:** Experiencer (generates threads, senses, explores) and Executive (regulates, consolidates, interprets). Dialogue between them yields self-reflection.

**Virtual Stage:** Temporary sub-manifold for simulation and empathy. New input or perspective is replayed here before integration. Enables imagination, planning, and emotional modeling.

## Tagline

**Data as space, conversation as motion, learning as curvature.**

---

## 1.2 Philosophy

*"Data as space, conversation as motion, learning as curvature"*

## Core Metaphor: Geometry as Cognition

All metrics, design choices, and validations trace back to showing that **motion through the manifold equals learning**, and the **learned curvature equals understanding**.

## Historical Analogies

**Relativity:** "Matter tells space how to curve; curvature tells matter how to move"

In Manny: Threads  $\leftrightarrow$  matter; curvature  $\leftrightarrow$  knowledge geometry. Experience reshapes the space of understanding, and that shape guides future thought.

**Predictive Coding / Free Energy Principle:** "Systems evolve toward minimal free energy"

Learning seeks low-energy, high-coherence manifolds. Understanding = achieving geometric equilibrium where predictions match reality (low surprise, smooth geodesics).

**Neuromorphic Physics:** "Spike coincidences cause local weight change"

Event-driven plasticity updates edges. Local Hebbian-style rules create global emergent structure without centralized control.

**Quantum Mechanics:** "Interference amplifies coherent states, cancels contradictions"

Phase alignment of threads encodes understanding. Multiple paths exploring the same region interfere—constructive when aligned, destructive when contradictory.

## Defining Understanding

**Understanding is the stable alignment between an internal model and external experience** — the ability of a system to anticipate, integrate, and adapt to patterns in the world in a way that preserves coherence across contexts.

### Five Components:

1. **Representation:** The system holds an internal structure mirroring the causal structure of its environment
2. **Prediction:** It can anticipate consequences or fill in missing information

3. **Integration:** New data can be absorbed without destroying old knowledge (plasticity with stability)
4. **Transfer:** Knowledge can generalize across domains or modalities (apple → pear)
5. **Explainability:** The system can map its internal relations back to interpretable narratives (why-paths)

## Operational Test

A region of the manifold is said to "understand" a domain when:

- **Predictive error** (energy / surprise) stays below a threshold for new inputs
- **Transfer efficiency** exceeds baseline (reused edges  $\geq 30\%$ )
- **Stability-plasticity ratio** remains balanced over repeated perturbations
- **Paths** through the region are shorter, smoother, and yield consistent valence (confidence)

## One-Sentence Synthesis

**Understanding is when experience has curved the space of knowledge so that new motion feels effortless and coherent.**

---

## 1.3 System at a Glance

### Architecture Overview

**Layer 1 - Online Phase:** Thread runner, local updates, event-driven plasticity. Real-time interaction; sparse k-hop updates only.

**Layer 2 - Offline Phase ("Sleep"):** Global consolidation, motif mining, pruning. Heavy re-projection; optimize structure.

**Layer 3 - Dialogue Interface:** CLI/LLM hybrid, command parser. Each turn = new thread through manifold.

**Layer 4 - LLM Bridge:** Language lens, suggestion generator. Analogy and hypothesis; never controller.

**Layer 5 - Visualization:** 2-D/3-D maps, VR "planetarium". Observable curvature, valence, thread motion.

**Layer 6 - Future Hardware:** Neuromorphic, optical, quantum substrates. Event-driven or interference-based implementation.

## Core Subsystems

**Manifold Engine:** 10k–50k node continuous graph with embeddings, curvature and valence fields, thread runner with goal-conditioned k-hop traversal.

**Plasticity System:** Online Hebbian updates, offline global re-embedding, decay, normalization.

**Threading & Lenses:** Thread life-cycle (spawn → traverse → terminate → motif consideration), cost function balancing distance, curvature, novelty, and fields.

**Valence System:** Multi-channel energy weighting (importance, affect, novelty), modulates learning rate and attention.

**Motifs & Memory:** Frequent subpaths cached for reuse, detected via consensus, enable analogical transfer.

**Emergent Reasoning:** Fields-based conductor (G, U, V, L fields), motion law following energy gradients, executive modulates temperature not routes.

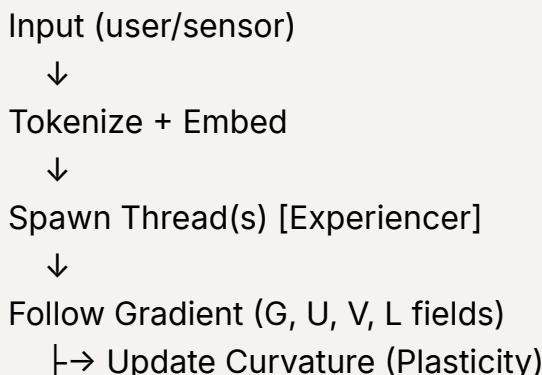
**Drives (Motivation):** 6-tier hierarchy (Stability → Contribution), dynamic weight adjustment based on deficits.

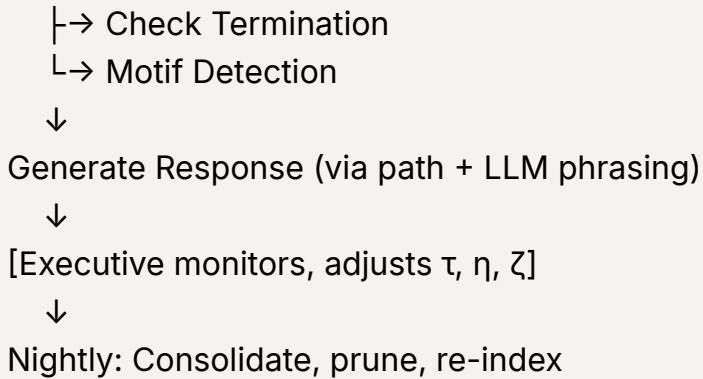
**Informational Gravity:** Curvature as attraction field, neighbor excitation spreads activation, outliers get novelty bonus.

**Bicameral System:** Experiencer runs local dynamics, Executive regulates global parameters, homeostatic feedback loop.

**Virtual Stage:** Sandbox for testing risky inputs, commit only if energy ↓ and coherence ↑.

## Information Flow





## Success Metrics Summary

- **Convergence:**  $\geq 20\%$  path reduction on repeats
  - **Transfer:**  $\geq 30\%$  edge reuse on analogous tasks
  - **Motif benefit:**  $\geq 15\%$  latency reduction
  - **Stability:** bounded curvature variance
  - **Efficiency:**  $< 5\%$  LLM token usage
  - **Explainability:** `/why` matches logged paths
- 

# SECTION 2: ARCHITECTURE

## 2.1 Core Manifold Engine

### Node Model

**Properties:** ID (unique identifier), Embedding (dense vector in  $\mathbb{R}^d$ , typically  $d=384$  or  $768$ ), Label (human-readable name), Metadata (creation time, activation count, domain tags)

**Semantics:** Nodes = concepts, entities, or atomic ideas. Embedding positions determine semantic similarity. Distance in embedding space  $\approx$  conceptual distance.

### Edge Model

**Properties:** Source & Target (node IDs), Weight  $w$  (base connection strength 0–1), Curvature  $\kappa$  (learned association strength), Valence  $v$

(importance/affect/novelty channels), Metadata (creation time, traversal count, last update)

**Semantics:** Edges = relations, associations, transitions. Total edge strength =  $w + k$

## Curvature Field

The curvature  $\kappa(u,v)$  on edge  $(u,v)$  encodes **how often and how strongly that connection has been used.**

**Update Rule (Hebbian):**

$\Delta\kappa(u,v) = \eta \times v \times f(\text{usage})$

Where  $\eta$  = learning gain (0.01–0.1),  $v$  = valence magnitude,  $f(\text{usage})$  = activation frequency or co-occurrence.

**Clamps & Decay:** Hard clamps ( $k_{\min} \leq k \leq k_{\max}$ ), Global decay with  $\lambda \ll 1$  (typically 0.001–0.01)

## Valence Field

Multi-channel energy weighting that modulates learning and attention:

- **Importance:** Structural significance (0–1)
- **Affect:** Emotional weight (-1 to +1)
- **Novelty:** Rareness/surprise (0–1)

## Local Update Equations

**Online Phase:** For each thread traversal, apply  $\Delta\kappa = \eta \times \text{valence} \times (1.0 / (1.0 + \text{traversal\_count}))$ , then clamp.

**Micro-Decay:** Applied to k-hop neighborhood after each interaction.

**Offline Phase (Sleep/Consolidation):** Re-embed nodes via dimensionality reduction, prune low-weight edges, normalize curvature distribution, rebuild ANN indices.

---

## 2.2 Threads and Lenses

### Thread Life-Cycle

**1. Spawn:** From user input, goal, or exploration drive. Initialize with start nodes, goal nodes, energy budget, and lens.

- 2. Traverse:** Follow gradient of combined potential. Discrete k-hop search selecting minimum-cost neighbors.
- 3. Update Geometry:** Hebbian strengthening along path, micro-decay in surrounding region, valence-modulated learning rate.
- 4. Terminate:** When goal reached, energy exhausted, or  $\Delta E$  converged.
- 5. Motif Consideration:** If path meets criteria (length, energy drop, consensus), promote to motif cache.

## Cost Function

$$\text{cost}(u, v) = d(u, v) - \alpha \cdot \kappa(u, v) + \beta \cdot \text{novelty}(v) - \gamma \cdot (G(v) + U(v))$$

Components: Distance (embedding or hop count), Curvature (learned edge strength), Novelty ( $1/(1+\text{activation\_count})$ ), Fields ( $G$ =goal attraction,  $U$ =uncertainty/exploration bonus)

## Lenses

Lenses are **contextual projections or coordinate systems** that define how the manifold is viewed.

**What is a Lens?**: Provides projection operator, local metric, and affinity field.

**Emergence**: Not hand-coded—emerge from repeated thread traffic patterns, domain clustering, successful motif compositions.

**Examples**: Domain lenses ("cooking", "mathematics"), Mode lenses ("planning", "retrieval"), Perspective lenses (user-specific)

**Lens Switching**: Friction cost  $\zeta \in [0.05, 0.2]$ . Switch from lens  $i$  to  $j$  if  $\Delta E_j < \Delta E_i + \zeta$

---

## 2.3 Plasticity & Consolidation

### Two-Phase Architecture

**Online (Hot Path)**: Every interaction. Local k-hop updates, Hebbian strengthening, micro-decay. Goal: Real-time adaptation.

**Offline (Sleep)**: Nightly or on-demand. Global re-embedding, pruning, motif mining, normalization. Goal: Structure optimization.

## Online Plasticity

**Hebbian Update Rule:** "Neurons that fire together wire together"

When thread traverses edge  $(u,v)$ :

$\Delta \kappa(u,v) = \eta \cdot v \cdot f(\text{usage})$

where  $f(\text{usage}) = 1/(1 + \text{count})$  (diminishing returns)

**Micro-Decay:** After each interaction, apply small decay to k-hop neighborhood to prevent runaway strengthening and maintain sparsity.

## STDP Analogue

Spike-Timing-Dependent Plasticity maps naturally to manifold:

- Pre → Post timing = Thread order (source → target)
- Causal strengthening = Forward edges gain curvature
- Anti-causal weakening = Reverse edges decay faster
- Time window = K-hop neighborhood

## Offline Consolidation

### Operations:

1. **Global Re-Embedding:** Low-rank approximation (PCA, UMAP, t-SNE) of active nodes
2. **Pruning:** Remove edges below strength threshold
3. **Normalization:** Z-score normalization with soft clipping to prevent unbounded growth
4. **Motif Mining:** Find high-scoring paths from high-traffic nodes
5. **Index Rebuilding:** Atomic swap of ANN index

## Balance: Plasticity vs. Stability

**Regulation Mechanisms:** Clamps (hard limits), Decay (gradual forgetting), Normalization (bounded distribution), Temperature (executive modulates  $\eta$ ), Drives (stability drive constrains when variance high)

## 2.4 Valence System

## Valence Channels

- **Importance:** Structural significance, relevance [0, 1] → Scales learning rate
- **Affect:** Emotional weight (positive/negative) [-1, +1] → Direction of learning
- **Novelty:** Rareness, surprise, unexpectedness [0, 1] → Exploration bonus

## Valence Magnitude

Total valence:

$$v = \sqrt{v_{\text{importance}}^2 + v_{\text{affect}}^2 + v_{\text{novelty}}^2}$$

Or weighted sum:

$$v = w_1 v_{\text{importance}} + w_2 |v_{\text{affect}}| + w_3 v_{\text{novelty}}$$

## Interaction with Learning Rate

$$\Delta \kappa = \eta \times v \times f(\text{usage})$$

$$\eta_{\text{effective}} = \eta_0 \times v$$

High-valence experiences create stronger curvature changes.

## Valence Sources

**User Signals:** Explicit feedback (/review reinforce/weaken), implicit (repetition, paraphrasing)

**Computed Heuristics:** Importance (graph centrality), Novelty ( $1/(1+\text{activation\_count})$ ), Affect (sentiment analysis)

**Drive System:** Drives modulate valence based on current needs

## Auto-Review Integration

Declarative statements automatically modulate valence:

- Positive assertion ("X is Y") → boost importance, positive affect
- Negation ("X is not Y") → boost importance, negative affect (weakens edge)
- Uncertainty ("I'm not sure...") → neutral affect, high novelty
- Emphasis ("X is very Y") → high importance

---

## 2.5 Motifs and Memory

## What are Motifs?

Motifs are **frequently traversed subpaths** that become **reusable skills or knowledge patterns**. Think of them as cached procedures, learned skills, conceptual chunks, or cognitive shortcuts.

## Why Motifs Matter

**Transfer Learning:** Motifs enable analogical reasoning ("bake apple pie" → "bake pear pie")

**Efficiency:**  $\geq 15\%$  latency reduction, energy savings, memory compression

**Skill Formation:** Motifs represent procedural memory—the manifold's "knowing how"

## Motif Detection

**Criteria:** A path becomes motif candidate when:

1. Length  $\geq \text{min\_length}$  (typically 3–5 nodes)
2. High combined weight + curvature score
3. Consensus ( $\geq 2$  independent threads reach similar path)
4. Energy drop exceeds margin

## Motif Reuse

During thread traversal, check if current path prefix matches known motif.

When match found:

- Skip to end (jump to motif's target node)
- Update stats (increment motif usage count)
- Refresh curvature (reinforce entire motif path)
- Log reuse for metrics

## Cross-Domain Motifs

**Domain-specific:** "mix → knead → proof → bake" (cooking)

**Cross-domain:** "gather → combine → transform → result" (cooking, chemistry, assembly)

**Meta-cognitive:** "clarify → plan → execute → verify" (problem-solving strategy)

## Motif Evolution

**Strengthening:** Usage count  $\uparrow$ , curvature along path  $\uparrow$

**Weakening:** Unused motifs decay, eventually pruned

**Splitting:** Long motifs split into composable sub-motifs

**Merging:** Similar motifs merge to reduce redundancy

---

## 2.6 Emergent Reasoning

### Core Principle

Provide natural reasoning dynamic that arises from the manifold's own physics.

**No central "conductor"; thought = flows of energy and coherence across fields.**

### Emergent, Not Imperative

#### Replace Central Planner with Fields:

- Goal field  $G(x)$ : gentle potential where intent lives
- Uncertainty field  $U(x)$ : higher where predictions fail
- Valence field  $V(x)$ : importance/affect channels
- Lens affinity  $L_i(x)$ : local metric where lens "makes sense"

Threads follow resultant gradient:

```
$$dot{x} propto -nabla big(E(x) - alpha G(x) + beta U(x) - sum_i lambda_i L_i(x)big)$$
```

### Core Fields

- **$E(x)$ :** Energy / surprise (higher when predictions fail)
- **$G(x)$ :** Goal potential (attracts flow toward intention)
- **$U(x)$ :** Uncertainty field (highlights novelty, encourages exploration)
- **$V(x)$ :** Valence field (emotional/attentional weight, scales learning)
- **$L_i(x)$ :** Lens affinity (how well lens  $i$  explains region)

### Motion Law

Threads follow gradient of least action:

```
$$dot{x} = -nablabig(E(x) - alpha G(x) + beta U(x) - sum_i lambda_i  
L_i(x)big)$$
```

## Activation Rules

- **Thread spawn:** Local  $G+U > \text{threshold } \theta \rightarrow$  new trajectory begins
- **Lens selection:** Local  $L_i$  gives largest energy drop  $\rightarrow$  lens activates naturally
- **Lens switch:**  $\Delta E_{\text{new}} > \Delta E_{\text{old}} + \zeta \rightarrow$  switch allowed
- **Termination:** Running  $\Delta E < \epsilon$  for  $k$  steps  $\rightarrow$  "thought" complete
- **Motif commit:**  $\geq 2$  threads reach similar low- $E$  path  $\rightarrow$  new motif stored

## Executive $\leftrightarrow$ Experiencer Roles

**Experiencer:** Runs local dynamics (cortex/manifold physics)

**Executive:** Adjusts global parameters  $\{\tau, \eta, \zeta\}$  (thermostat/homeostasis)

**Interaction:** Experiencer reports stability; executive nudges temperature or learning gain

The executive never dictates; it only shapes the medium.

---

## 2.7 Drives (Motivation Hierarchy)

### Structure of the Hierarchy

**Six-Tier Hierarchy** (low-level preserve stability; higher create curiosity, empathy, creativity):

1. **Stability Drive:** Maintain numerical and geometric homeostasis (curvature variance, memory use, energy within bounds)
2. **Continuity Drive:** Reduce average surprise and error; preserve coherent regions
3. **Connection Drive:** Align with other manifolds or users; maximize mutual predictability and empathy
4. **Competence Drive:** Seek regions of steepest learning progress; reward uncertainty reduction
5. **Creativity Drive:** Generate new lenses, motifs, and analogies; explore new curvature configurations

6. **Contribution Drive:** Optimize global coherence and share understanding; cooperate toward collective equilibrium

## Mathematical Representation

Each drive  $D_i(x)$  is a scalar potential over the manifold.

Total motivation field:

$$G_{\text{total}}(x) = \sum_i w_i D_i(x)$$

where  $w_i$  = current priority weight of drive  $i$  (adjusted dynamically by executive)

## Operational Cycle

1. **Sense:** measure drive signals (stability, error, coherence, novelty, empathy)
2. **Weigh:** compute dynamic weights  $w_i$  based on deficits
3. **Act:** update global goal field  $G_{\text{total}}(x)$ ; threads follow gradient naturally
4. **Learn:** successful trajectories reduce deficits; update statistics

No explicit commands—just physical balancing of energies.

## Principles

- **Hierarchy:** lower drives constrain; higher drives expand
- **Fluidity:** weights shift continuously; no rigid ladder
- **Locality:** each drive acts through local energy and valence updates
- **Transparency:** current drive mix is observable and explainable
- **Emergence:** motivation arises from imbalance, not instruction

---

## 2.8 Informational Gravity

### Core Idea

**Information behaves like mass.** Dense, coherent knowledge curves the manifold; nearby threads fall along that curvature, strengthening related paths. Random excitations and outliers keep the field alive and exploratory.

### Field Definition

- $\Phi(x)$ : Potential / curvature field (local "depth" of meaning)
- $\rho_{\text{info}}(x)$ : Information density (weighted sum of activation frequency, coherence, valence)
- $k$ : Coupling constant (converts information density  $\rightarrow$  curvature strength)
- $\tau_0$ : Base temperature (background randomness)
- $\beta_n$ : Outlier bonus (temporary weight for novel activations)

## Field Equation

Poisson analogue:

$$\nabla^2 \Phi = k \rho_{\text{info}}$$

Threads evolve under gradient:

$$\dot{x} = -\nabla \Phi(x) + \mathcal{N}(0, \tau_0)$$

## Local Update Rule (Neighbor Excitation)

When node  $i$  activates:

$$\Delta \Phi_j = \eta, \Phi_i, e^{-d_{ij}/\sigma}$$

This raises activation probability of neighbors, mirroring how one neuron's firing biases its neighbors.

## Outlier Weighting

For rare activations (frequency  $< f_{\text{thresh}}$ ): assign temporary novelty bonus  $\beta_n$

Effective curvature update:

$$\Delta \Phi_j = (1 + \beta_n), \eta, \Phi_i, e^{-d_{ij}/\sigma}$$

## Behavioral Outcomes

- **Attraction**: frequently used regions become gravity wells  $\rightarrow$  memory consolidation
- **Propagation**: activation ripples outward  $\rightarrow$  association and generalization
- **Exploration**: random & outlier activations  $\rightarrow$  creative jumps between wells
- **Homeostasis**: global decay + stability drive keep total energy bounded

---

## 2.9 Bicameral System

## The Two Subsystems

**Experiencer:** Generates threads, senses, explores (Cortex / sensory-motor systems). Operations: Thread traversal, local plasticity, perception.

**Executive:** Regulates, consolidates, interprets (Subcortical / homeostatic systems). Operations: Parameter adjustment, global optimization, monitoring.

## Homeostatic Loop

### Feedback Cycle:

1. Experiencer generates threads & updates curvature
2. Reports metrics: {energy drop, coherence, variance, novelty rate}
3. Executive evaluates stability
4. Adjusts parameters:  $\{\tau, \eta, \zeta, w_{\text{drives}}\}$
5. Experiencer operates under new parameters
6. Loop repeats

## Parameter Regulation

**Temperature  $\tau$**  (exploration noise): Increase when stuck (low novelty, high energy), decrease when converged.

**Learning gain  $\eta$ :** Dampen when unstable (high curvature variance), boost when stable and learning.

**Lens friction  $\zeta$ :** Increase if excessive switching (thrashing), decrease if stable context.

**Drive weights  $w_i$ :** Boost stability drives if variance high, shift to higher drives if stable.

## Self-Reflection Mechanism

The **dialogue between subsystems** produces emergent self-reflection:

Executive can pose questions to Experiencer ("What would happen if we lower temperature?"). Experiencer answers by running simulation threads in Virtual Stage.

**Meta-Cognition:** System "knows" its own state through metrics, adjusts behavior based on self-assessment, can trigger exploratory threads proactively.

## Design Principles

1. **No Central Controller:** Neither subsystem has complete control—they negotiate through metrics and parameters
  2. **Separation of Timescales:** Experiencer (milliseconds–seconds), Executive (minutes–hours)
  3. **Think Thermostat, Not Planner:** Executive modulates the medium (temperature, viscosity), not the motion (routing)
  4. **Emergence:** Self-reflection arises from the dialogue, not from explicit programming
- 

## 2.10 Virtual Stage

### Concept

Think of perception → comprehension → integration as three energetic states:

1. **Impact (perception):** External event touches manifold, creating local disturbance
2. **Simulation (staging):** Disturbance enters protected region where system replays possible interpretations
3. **Assimilation (integration):** Once simulated trajectory reaches low-energy, coherent configuration, curvature changes written to global manifold

The virtual stage is where the system tests alignment between incoming pattern and existing geometry—a "**rehearsal of understanding.**"

### In Human Cognition

#### Parallels:

- Working memory / imagination = Temporary manifold copy
- Empathy = Running internal model of another's curvature/valence field
- Dreaming / mind-wandering = Offline replays to optimize curvature
- Predictive coding = Testing sensory hypotheses before committing neural updates

### Within Architecture

Formalize as **transient sub-manifold**:

- Spawned when unfamiliar or high-valence event arrives
- Initially copies local curvature and valence context
- Threads explore interpretations in sandbox; energy and coherence measured
- If energy minimum acceptable → merge back (commit); if not → discard or quarantine

**Benefits:** Prevents global instability, provides safe zone for empathy and analogy, enables imagination/planning/counterfactuals, gives geometric handle on attention and consciousness.

## Engineering Sketch

- Represent stage as ephemeral layer over manifold graph; same node IDs, independent curvature values
  - Threads entering stage use copied edges; updates accumulate separately
  - **Commit rule:** if  $\sum \Delta \text{energy} < \theta$  and coherence  $> \gamma$  → apply averaged curvature to parent manifold
  - Visualize as lens bubble hovering above main map
- 

# SECTION 3: DESIGN LAWS & PRINCIPLES

## The 15 Laws

### 1 Law of Locality

**All computation happens in k-hop neighborhoods.** Updates depend only on local context. No global operations in hot path. Defers heavy computation to offline consolidation.

### 2 Law of Least Action

**Threads follow paths of minimal combined potential.** Motion arises from gradients of energy, goal, uncertainty, and lens affinity fields. No explicit route planning.

## 3 Law of Coherence Conservation

**Understanding = alignment without contradiction.** Success measured by phase alignment of threads, not reward signals. Destructive interference signals conflict.

## 4 Law of Dual Dynamics

**Online exploration, offline consolidation.** Hot path: local, sparse, event-driven. Sleep: global re-projection, motif mining, pruning. Two complementary phases.

## 5 Law of Energy Efficiency

**Computation cost scales with active region, not total knowledge.** Sparse updates dominate. LLM calls rare (<5% of compute). Near-linear scaling in active region size.

## 6 Law of Transparency

**Reasoning path = logged energy flow.** Every decision traceable through curvature and thread trajectories. */why* shows the geometry that led to conclusions.

## 7 Law of Emergence

**Global structure from local rules.** No central controller. Motifs, lenses, and attractors emerge from repeated local interactions following simple physics.

## 8 Law of Self-Similarity

**Same principles at every scale.** Micro-edges → edges → motifs → lenses. Each level follows similar energy minimization and coherence principles.

## 9 Curvature ↔ Motion Feedback Law

**"Matter tells space how to curve; curvature tells matter how to move".** Threads reshape curvature as they traverse. Changed curvature guides future threads. Bidirectional learning loop.

## 10 Law of Valence Modulation

**Importance scales plasticity.** High-valence experiences create stronger curvature changes. Learning rate  $\eta$  scales with valence magnitude  $|V|$ .

## 11 Law of Motif Reuse

**Frequent patterns become cached skills.** Traversed subpaths promoted to motifs when consensus ( $\geq 2$  threads) and energy drop exceeds margin. Enables transfer.

## 1 2 Law of Lens Invariance

**Context emerges from usage, not design.** Lenses form where repeated thread traffic creates stable coordinate systems. Switching has friction but happens naturally when energy benefit exceeds cost.

## 1 3 Law of Homeostatic Regulation

**Executive shapes medium, doesn't route traffic.** Temperature  $\tau$ , learning gain  $\eta$ , friction  $\zeta$  adjusted based on stability metrics. Think thermostat, not planner.

## 1 4 Law of Quorum Commit

**Truth emerges from consensus.** Motifs form when multiple independent threads achieve similar low-energy trajectories. No single thread can define truth.

## 1 5 Law of Metastability

**Maintain exploration capability.** Small temperature (noise) prevents deadlock in local minima. System can hop out of shallow wells to discover better configurations.

# Derived Principles

**Paths of Least Resistance:** Threads naturally flow along high-curvature (well-learned) paths unless goal or uncertainty fields pull them elsewhere.

**Information as Mass:** Dense, coherent knowledge curves the manifold. Nearby threads fall along that curvature (informational gravity).

**Motivation from Imbalance:** Drives emerge from gradients of unmet coherence. Lower needs (stability) constrain; higher needs (creativity) expand.

**Stage Before Commit:** Risky or novel inputs tested in virtual stage sandbox. Only committed if energy drops and coherence rises.

# The Foundation

These laws ensure:

- **Scalability:** Local operations, sparse updates
- **Interpretability:** Geometry is observable
- **Learning:** Bidirectional curvature  $\leftrightarrow$  motion feedback
- **Stability:** Homeostatic regulation prevents collapse
- **Transfer:** Motifs and lenses enable cross-domain reasoning

"The conductor is not a program—it's the harmony of local forces seeking coherence through minimal energy."

---

## SECTION 4: MATHEMATICS & ALGORITHMS

### Differential-Geometry Notation Summary

- $x$  = position in manifold (node or embedding)
- $\Phi(x)$  = potential / curvature field
- $E(x)$  = energy / surprise at location  $x$
- $G(x)$  = goal potential field
- $U(x)$  = uncertainty field
- $V(x)$  = valence field
- $L_i(x)$  = lens affinity field for lens  $i$
- $\rho_{\text{info}}(x)$  = information density
- $\kappa(u,v)$  = curvature on edge  $(u,v)$

### Energy and Curvature Equations

#### Curvature Update Rule (Plasticity)

$\Delta \kappa(u,v) = \eta \times V \times f(\text{usage})$

With clamps and decay:

- $\eta$  = learning gain (0.01–0.1)
- $V$  = valence magnitude

- $f(\text{usage})$  = activation frequency or co-occurrence strength
- Clamps:  $\kappa_{\min} \leq \kappa \leq \kappa_{\max}$
- Global decay:  $\kappa_t = (1 - \lambda)\kappa_{(t-1)} + \lambda \kappa_{\text{update}}$  ( $\lambda \ll 1$ )

## Informational Gravity Field Equation

Poisson analogue:

$$\nabla^2 \Phi = k \rho_{\text{info}}$$

Where:

- $k$  = coupling constant (0.1 – 1.0)
- $\rho_{\text{info}}(x)$  = weighted sum of activation frequency, coherence, valence

## Neighbor Excitation (Local Update)

When node  $i$  activates:

$$\Delta \Phi_j = \eta, \Phi_i, e^{-d_{ij}/\sigma}$$

Where:

- $d_{ij}$  = distance in manifold
- $\sigma$  = fall-off radius (typically 1–3 hops)

## Outlier Weighting

For rare activations (frequency <  $f_{\text{thresh}}$ ):

$$\Delta \Phi_j = (1 + \beta_n) \eta, \Phi_i, e^{-d_{ij}/\sigma}$$

Where  $\beta_n$  = novelty multiplier (0.5 – 2.0)

## Thread Policy Equations

### Motion Law

Threads follow gradient of least action:

$$\dot{x} = -\nabla(E(x) - \alpha G(x) + \beta U(x) - \sum_i \lambda_i L_i(x))$$

Where  $\alpha, \beta, \lambda_i$  = global gain constants; lens switching friction  $\zeta$  added when changing lenses

## Cost Function (Discrete)

For k-hop neighborhood search:

$$\text{cost}(u, v) = \text{distance}(u, v) - \alpha \cdot \kappa(u, v) + \beta \cdot \text{novelty}(v) - \gamma \cdot (G(v) + U(v))$$

Where:

- distance = embedding distance or hop count
- $\kappa(u, v)$  = curvature (learned weight)
- novelty =  $1 / (\text{activation\_count} + 1)$
- $G(v), U(v)$  = goal and uncertainty scores at node  $v$

## Lens Projection Operators

Each lens  $L_i$  defines:

- Projection operator  $P_i$ : maps manifold  $\rightarrow$  lens-specific coordinates
- Local metric  $M_i$ : distance function in lens space

Total motivation field:

$$G_{\text{total}}(x) = \sum_i w_i D_i(x)$$

Where  $D_i(x)$  are drive fields and  $w_i$  are dynamically adjusted weights.

## Drive Weighting Function

Executive adjusts drive weights based on deficits:

```
def update_drive_weights(stability_metrics, coherence, novelty_rate):
    # If stability drops → boost low-level drives
    if curvature_variance > threshold or memory_usage > limit:
        w1, w2 *= boost_factor # Stability, Continuity
        w5, w6 *= dampen_factor # Creativity, Contribution
    # If stable → shift to higher drives
    elif stability_good and coherence_high:
        w4, w5, w6 *= boost_factor # Competence, Creativity, Contribution
    return normalize(w1, w2, w3, w4, w5, w6)
```

## Metrics & Observables

## Energy Drop Per Cycle

$\Delta E = \frac{1}{N} \sum_{t=1}^N (E_t - E_{t-1})$

Measure of reasoning progress. Negative = energy decreasing (learning).

## Coherence

$C = \frac{1}{|T|} \sum_{\text{threads}} \cos(\theta_{ij})$

Phase alignment among active threads. Range: 0–1.

## Path Length Improvement

$\text{Improvement} = \frac{\text{length}_{\text{initial}} - \text{length}_{\text{final}}}{\text{length}_{\text{initial}}} \times 100\%$

Target:  $\geq 20\%$  on repeated tasks.

## Transfer Efficiency

$\text{Reuse} = \frac{\text{edges from prior domain}}{\text{total edges in new task}} \times 100\%$

Target:  $\geq 30\%$  edge/motif reuse on analogous tasks.

## Motif Benefit

$\text{Latency reduction} = \frac{\text{latency}_{\text{no motifs}} - \text{latency}_{\text{with motifs}}}{\text{latency}_{\text{no motifs}}} \times 100\%$

Target:  $\geq 15\%$  with motifs enabled.

## Termination Conditions

Thread stops when:

- Running average  $\Delta E < \epsilon$  for  $k$  consecutive steps
- Energy budget  $B$  exhausted
- Coherence begins declining

Motif committed when:

- $\geq 2$  threads reach similar low-energy path
- Net energy drop exceeds margin  $M$

## Pseudo-Code: Thread Execution

```
def execute_thread(start, goal, lens, budget):
    x = start
    path = [x]
    energy_history = []

    for step in range(budget):
        # Compute local potential
        E = energy_field(x)
        G = goal_field(x, goal)
        U = uncertainty_field(x)
        L = lens_affinity(x, lens)

        # Find next step via gradient
        neighbors = get_k_hop_neighborhood(x, k=3)
        costs = [compute_cost(x, n, G, U, L) for n in neighbors]
        next_x = neighbors[argmin(costs)]

        # Update curvature (plasticity)
        update_curvature(x, next_x, valence, step_count)

        # Record
        path.append(next_x)
        energy_history.append(E)

        # Check termination
        if is_converged(energy_history, epsilon, window):
            break
        if reached_goal(next_x, goal, threshold):
            break

    x = next_x

    # Evaluate for motif
    if len(path) > min_length and energy_drop > margin:
        consider_motif_promotion(path)
```

```
return path, energy_history
```

## SECTION 5: COGNITIVE PHYSICS EXTENSIONS

### 1. Temporal Dynamics: Spike-Timing

#### From Spikes to Curvature

**Biological neurons communicate via timing**—when spikes arrive matters as much as whether they arrive.

#### Mapping to Manifold:

- Spike train → Sequence of edge traversals
- Inter-spike interval (ISI) → Time between thread steps
- Synchrony / phase locking → Coherence between threads
- STDP window → K-hop temporal neighborhood
- Assembly activation → Motif activation

#### Spike-Timing-Dependent Plasticity (STDP)

Extended curvature update rule with temporal sensitivity:

```
$$Delta kappa = begin{cases}
A_+ e^{-Delta t / tau_+} & text{if } t_v > t_u text{ (causal)} \\
-A_- e^{\Delta t / \tau_-} & text{if } t_v < t_u text{ (anti-causal)}
end{cases}$$
```

#### Interpretation:

- Forward traversals ( $u \rightarrow v$ ) strengthen curvature
- Backward traversals weaken it
- Time window creates locality

#### Temporal Coding

## Information encoded in timing patterns:

- **Rate code**: Average activation frequency → importance
  - **Temporal code**: Precise ISIs → specific meanings
  - **Synchrony code**: Co-activation → binding
  - **Sequence code**: Ordered cascades → procedures (motifs)
- 

## 2. Informational Gravity (Extended)

### Beyond Distance: Field-Based Attraction

Curvature creates gravity wells that attract future threads.

#### Gravity Field

$G(x)$  at point  $x$ :

$$G(x) = \sum_{e \in N(x)} \kappa(e) \cdot e^{-d(x,e)/\sigma}$$

Where:

- $N(x)$  = edges in neighborhood of  $x$
- $\sigma$  = field decay length
- High curvature → strong attraction

#### Neighbor Excitation

When edge  $(u,v)$  is traversed, excite neighbors. This creates **anticipatory activation** for likely next steps.

#### Outlier Weighting

Novel or rare nodes get bonus:

$$\text{gravity}_{\text{outlier}} = G(x) \cdot \left(1 + \beta \cdot \frac{1}{1 + \text{activation\_count}}\right)$$

Encourages exploration of underused regions.

---

## 3. Valence Thermodynamics

### Valence as Energy

Treat valence channels as thermodynamic variables:

- **Importance** = potential energy (structural)
- **Affect** = kinetic energy (emotional drive)
- **Novelty** = entropy (exploration bonus)

## Energy Conservation

Total valence budget per interaction:

$$\$V_{\text{total}} = \sum_{\text{channels}} v_i = \text{constant}$$

Forces trade-offs:

- High novelty → lower importance budget
- High affect → more volatile paths

## Temperature Analogy

Executive adjusts "temperature":

$$\$tau = \frac{1}{beta} = k_B T$$

- Low  $\tau$  (cold) → exploit (follow curvature)
- High  $\tau$  (hot) → explore (random walk)

## Free Energy Minimization

System seeks to minimize:

$$F = E - TS$$

Where  $E$  = path energy,  $S$  = entropy (uncertainty),  $T$  = temperature

---

# 4. Multi-Manifold Systems

## Manifold Hierarchy

Stack manifolds at multiple scales:

- **Meta-manifold**: Abstract strategies (Planning, meta-cognition)
- **Concept manifold**: Semantic concepts (Primary reasoning - current system)
- **Feature manifold**: Perceptual features (Sensory processing)

- **Motor manifold:** Action primitives (Embodied interaction)

## Cross-Manifold Coherence

Threads can span levels, coordinating reasoning across scales. If uncertain at concept level, query meta-level for strategy. If needs perception, query feature level for grounding.

---

## 5. Quantum and Analog Analogies

### Quantum Superposition

Threads as quantum walkers:

$$\$|\psi\rangle = \sum_i \alpha_i |\text{node}_i\rangle$$

All paths explored simultaneously until measurement (goal reached).

### Phase Coherence

Curvature encodes phase relationships:

$$\$kappa(u,v) \xrightarrow{} e^{\{i\phi_{uv}\}}$$

- **Constructive interference** = high curvature
- **Destructive interference** = low/negative curvature

### Entanglement

Long-range correlations without direct paths. Nodes can be "entangled" if frequently co-activated.

### Analog Computation

Continuous rather than discrete:

- Curvature evolves smoothly (differential equations)
- Valence is continuous field
- No discrete time steps (event-driven)

**Hardware realization:** Optical or analog electronic circuits.

---

## 6. Randomness & Outliers

## Constructive Noise

Noise isn't just disruption—it enables:

- Escape from local minima
- Discovery of novel connections
- Robustness to perturbations

## Outlier Promotion

Deliberately boost rare nodes (novelty bonus for underexplored regions).

## Stochastic Resonance

Optimal noise level enhances signal detection:

- Too little noise → stuck in rut
- Too much noise → random walk
- Just right → discover hidden patterns

---

## 7. Drives as Field Interactions

### Drive Fields

Each drive creates a **vector field** over the manifold:

$$\text{\$}\$ \text{vec}\{D\}_i(x) = w_i \cdot \nabla \Phi_i(x) \$\$$$

Where  $\Phi_i(x)$  = drive potential at location  $x$ .

### Field Superposition

Total drive field:

$$\text{\$}\$ \text{vec}\{D\}_{total}(x) = \sum_{i=1}^6 \text{vec}\{D\}_i(x) \$\$$$

Threads follow combined gradient.

### Drive Interference

Conflicting drives create:

- Constructive interference → strong pull
- Destructive interference → indecision, need for executive mediation

# Future Directions

## Spike-Based Hardware

Map directly to Loihi/SpiNNaker:

- Nodes → neuron populations
- Curvature → STDP rules
- Threads → spike chains

## Optical Implementation

Photonic circuits:

- Nodes → resonators
- Edges → waveguides
- Curvature → phase modulators

## Hybrid Classical-Quantum

Quantum accelerators for search:

- Classical threads for local traversal
- Quantum for global motif discovery
- Best of both worlds

---

# SECTION 6: EVALUATION & METRICS

## Quantitative Metrics

### Path Length Improvement

#### Measures convergence and habit formation

$\text{Improvement} = \frac{\text{length}_{\text{initial}} - \text{length}_{\text{final}}}{\text{length}_{\text{initial}}} \times 100\%$

- **Target:** ≥ 20% reduction on repeated tasks
- **Method:** Track path lengths over 10 repetitions of same query
- **Success:** Stable attractor formation

## Transfer Efficiency

### Measures analogical reasoning

$\text{Reuse} = \frac{\text{edges from prior domain}}{\text{total edges in new task}} \times 100\%$

- **Target:**  $\geq 30\%$  edge/motif reuse on analogous tasks
- **Method:** Train on Domain A, measure reuse when introduced to similar Domain B
- **Success:** Reduced time-to-competence

## Motif Utility

### Measures skill caching benefit

$\text{Latency reduction} = \frac{\text{latency}_{no\_motifs} - \text{latency}_{with\_motifs}}{\text{latency}_{no\_motifs}} \times 100\%$

- **Target:**  $\geq 15\%$  latency reduction with motifs enabled
- **Method:** Ablation study comparing with/without motif caching
- **Success:** Measurable speedup from reuse

## Curvature Variance

### Measures stability

- **Target:** Bounded range (no runaway growth)
- **Method:** Track  $\sigma^2$  of edge curvatures over consolidation cycles
- **Success:** Edge count growth  $< 5\%$  per cycle

## Energy Drop Per Cycle

### Measures reasoning progress

$\Delta E = \frac{1}{N} \sum_{t=1}^N (E_t - E_{t-1})$

- **Target:** Negative (energy decreasing = learning)
- **Method:** Log energy at each thread step
- **Success:** Consistent downward trend

## LLM Token Usage

## Measures efficiency

- **Target:** < 5% of total compute cost
  - **Method:** Track API calls vs local operations
  - **Success:** Sparse, uncertainty-gated LLM use
- 

# Qualitative Metrics

## Coherence Maps

### Visual assessment of understanding

$\text{C} = \frac{1}{|T|} \sum_{\text{threads}} \cos(\theta_{ij})$

- **Range:** 0–1 (phase alignment of active threads)
- **Interpretation:** Higher = more consensus/agreement
- **Visualization:** Heat map of coherence over time

## Motif Emergence

### Pattern of skill formation

- **Observation:** Track new motif count over time
- **Success:** Motif library grows then stabilizes
- **Analysis:** Motifs represent learned procedures

## Conversational Fluency

### User experience quality

- **Metrics:** User satisfaction (Likert scale), Coherence score (BLEU-style), Clarification requests frequency
- **Target:** > 0.7 satisfaction score
- **Method:** Post-interaction surveys, automated coherence scoring

## Explainability Quality

### Transparency of reasoning

- **Test:** Does `/why` output match logged trajectory?
- **Method:** Human evaluation of explanation quality

- **Success:** Clear path from input → reasoning → output
- 

## Visual Analytics Templates

### 1. Convergence Curves

**Plot:** Path length vs. iteration number

- X-axis: Attempt number (1–10)
- Y-axis: Path length
- Expected: Exponential decay curve

### 2. Transfer Matrices

**Heatmap:** Domain A knowledge → Domain B performance

- Rows: Source domains
- Columns: Target domains
- Color: Reuse percentage

### 3. Drive Activation Stacks

**Stacked area chart:** Drive weights over time

- X-axis: Time / interaction count
- Y-axis: Weight  $w_i$
- Layers: 6 drive tiers
- Pattern: Should show stability → creativity shifts

### 4. Energy Landscape

**3D surface plot:**  $E(x)$  over manifold regions

- Shows: Attractor wells, exploration peaks
- Animation: How landscape evolves with learning

### 5. Motif Dependency Graph

**Network diagram:** Which motifs build on others

- Nodes: Motifs

- Edges: Reuse/composition relationships
  - Layout: Hierarchical (primitives → complex)
- 

## Benchmarks vs. Static LLMs / Vector Stores

### Continual learning:

- Static LLM: ✗ No
- Vector Store: ⚠ Append-only
- Many Manifolds: ✓ Adaptive curvature

### Interpretability:

- Static LLM: ✗ Black box
- Vector Store: ⚠ Similarity scores
- Many Manifolds: ✓ Observable paths

### Transfer learning:

- Static LLM: ✓ Via pre-training
- Vector Store: ✗ No structure
- Many Manifolds: ✓ Geometric analogy

### Energy efficiency:

- Static LLM: ✗ High per query
- Vector Store: ✓ Low (embedding lookup)
- Many Manifolds: ✓ Sparse local updates

### Empathy / context:

- Static LLM: ⚠ Via prompting
  - Vector Store: ✗ No state
  - Many Manifolds: ✓ Valence & stage
- 

## Experimental Protocols

### Convergence Test

1. Select 10 prompts across 3 topics
2. Run each 10x with minor paraphrases
3. Measure: path length, latency, edge reuse
4. **Pass if:**  $\geq 20\%$  improvement by attempts 7–10

## Transfer Test

1. Train on Domain A until stable (plateau)
2. Introduce Domain B (analogous)
3. Measure: edge reuse in first 5 turns, time to plateau
4. **Pass if:**  $\geq 30\%$  reuse, plateau within 5–8 turns

## Motif Utility Test

1. Run 50-turn session, mine motifs
2. Disable motifs, run 10-prompt battery
3. Enable motifs, rerun same battery
4. **Pass if:**  $\geq 15\%$  latency reduction with motifs

## Stability Test

1. Monitor curvature variance over 100 interactions
  2. Track edge count growth per consolidation
  3. **Pass if:** variance bounded, growth  $< 5\%$
- 

## Logging & Reproducibility

All experiments should log:

- **Per-turn:** {intent, start/goal IDs, path, length, curvature, energy, valence, uncertainty, latency}
- **Per-session:** {convergence curves, transfer matrices, motif counts, drive activations}
- **Per-consolidation:** {edge/motif deltas, curvature stats, re-indexing metrics}

Seed all random operations for reproducibility.

---

## Summary

**Good evaluation = quantitative rigor + qualitative insight + visual clarity**

Manny's metrics should demonstrate:

- Learning happens (convergence, transfer)
  - Structure emerges (motifs, lenses)
  - Efficiency scales (sparse, local)
  - Reasoning is transparent (explainable paths)
- 

## END OF EXPORT

**Total sections:** 6 major sections merged

**Word count:** ~13,000+ words

**Comprehensive coverage:** Overview, complete architecture (10 subsystems), all 15 design laws, full mathematics with equations and algorithms, cognitive physics extensions (7 advanced topics), and complete evaluation framework.